

## GRPO vs PPO

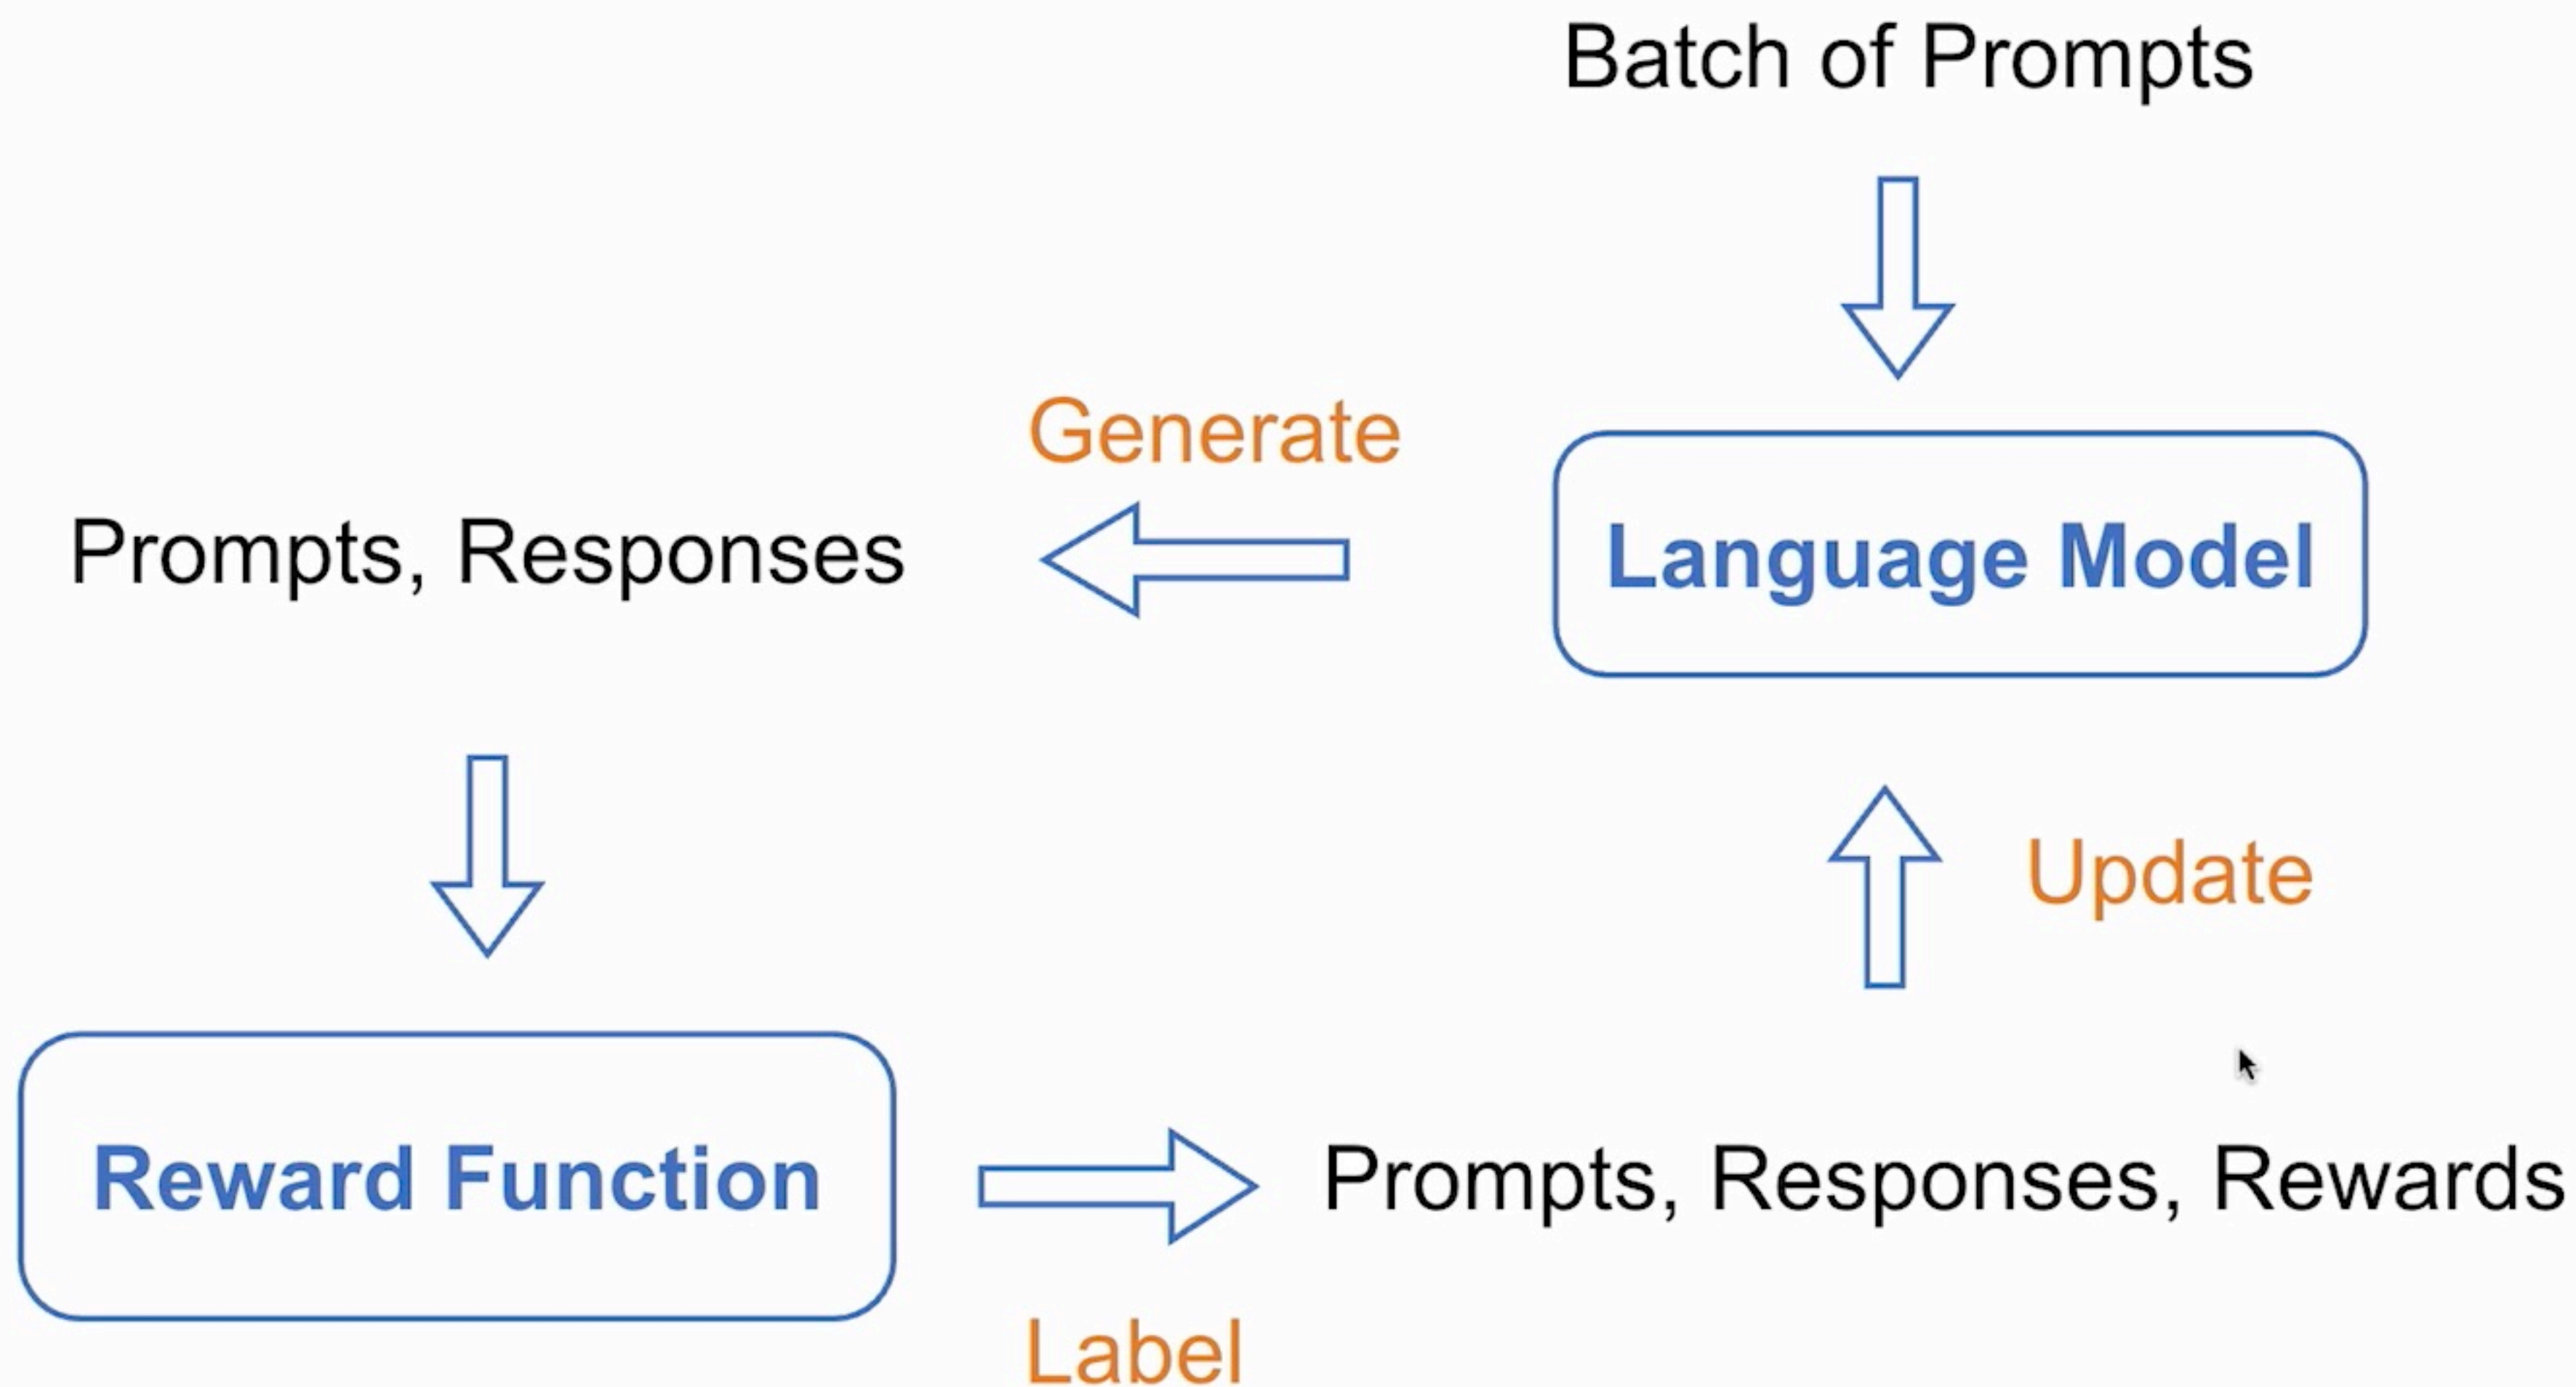
- Both GRPO and PPO are very effective online RL algorithms!
- **GRPO:**
  - Well-suited for binary (often correctness-based) reward
  - Requires larger amount of samples
  - Requires less GPU memory (no value model needed)
- **PPO:**
  - Works well with reward model or binary reward
  - More sample efficient with a well-trained value model
  - Requires more GPU memory (value model)

# Reinforcement Learning for LLMs: Online vs Offline

- **Online Learning:**
  - The model learns by generating new responses in real time — it iteratively collects new responses and their reward, updates its weights, and explores new responses as it learns.
- **Offline Learning:**
  - The model learns purely from a pre-collected prompt - response (-reward) tuple. No fresh responses generated during the learning process.



## Online RL: Let Model Explore Better Responses by Itself



# Reward Function in Online RL

## Option 1: Trained Reward Model

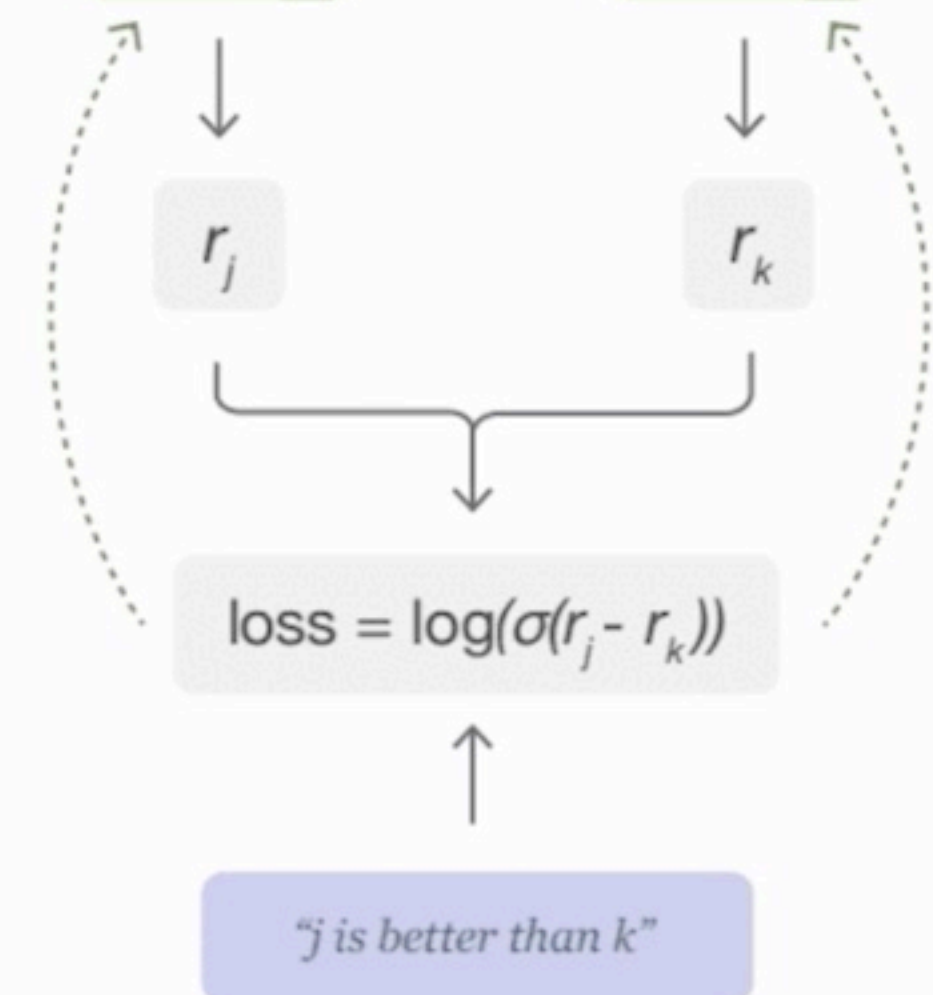
One post with two summaries judged by a human are fed to the reward model.



The reward model calculates a reward  $r$  for each summary.



The loss is calculated based on the rewards and human label, and is used to update the reward model.



- Usually initialized from an existing instruct model, then trained on large-scale human / machine generated preferences data
- Works for any open-ended generations;
- Good for improving chat & safety
- Less accurate for correctness-based domains like coding, math, function calling etc.




# Reward Function in Online RL

## Option 2: Verifiable Reward

**Math:** Check if the response matches ground truth

**Prompt:** What is  $1+1-1+1.1-1$

**Response:** The answer is  $\boxed{1.1}$ . 

**Ground truth:** 1.1

**Coding:** Running unit tests

**Prompt:** Given a string S, return the longest substring that occurs at least twice.

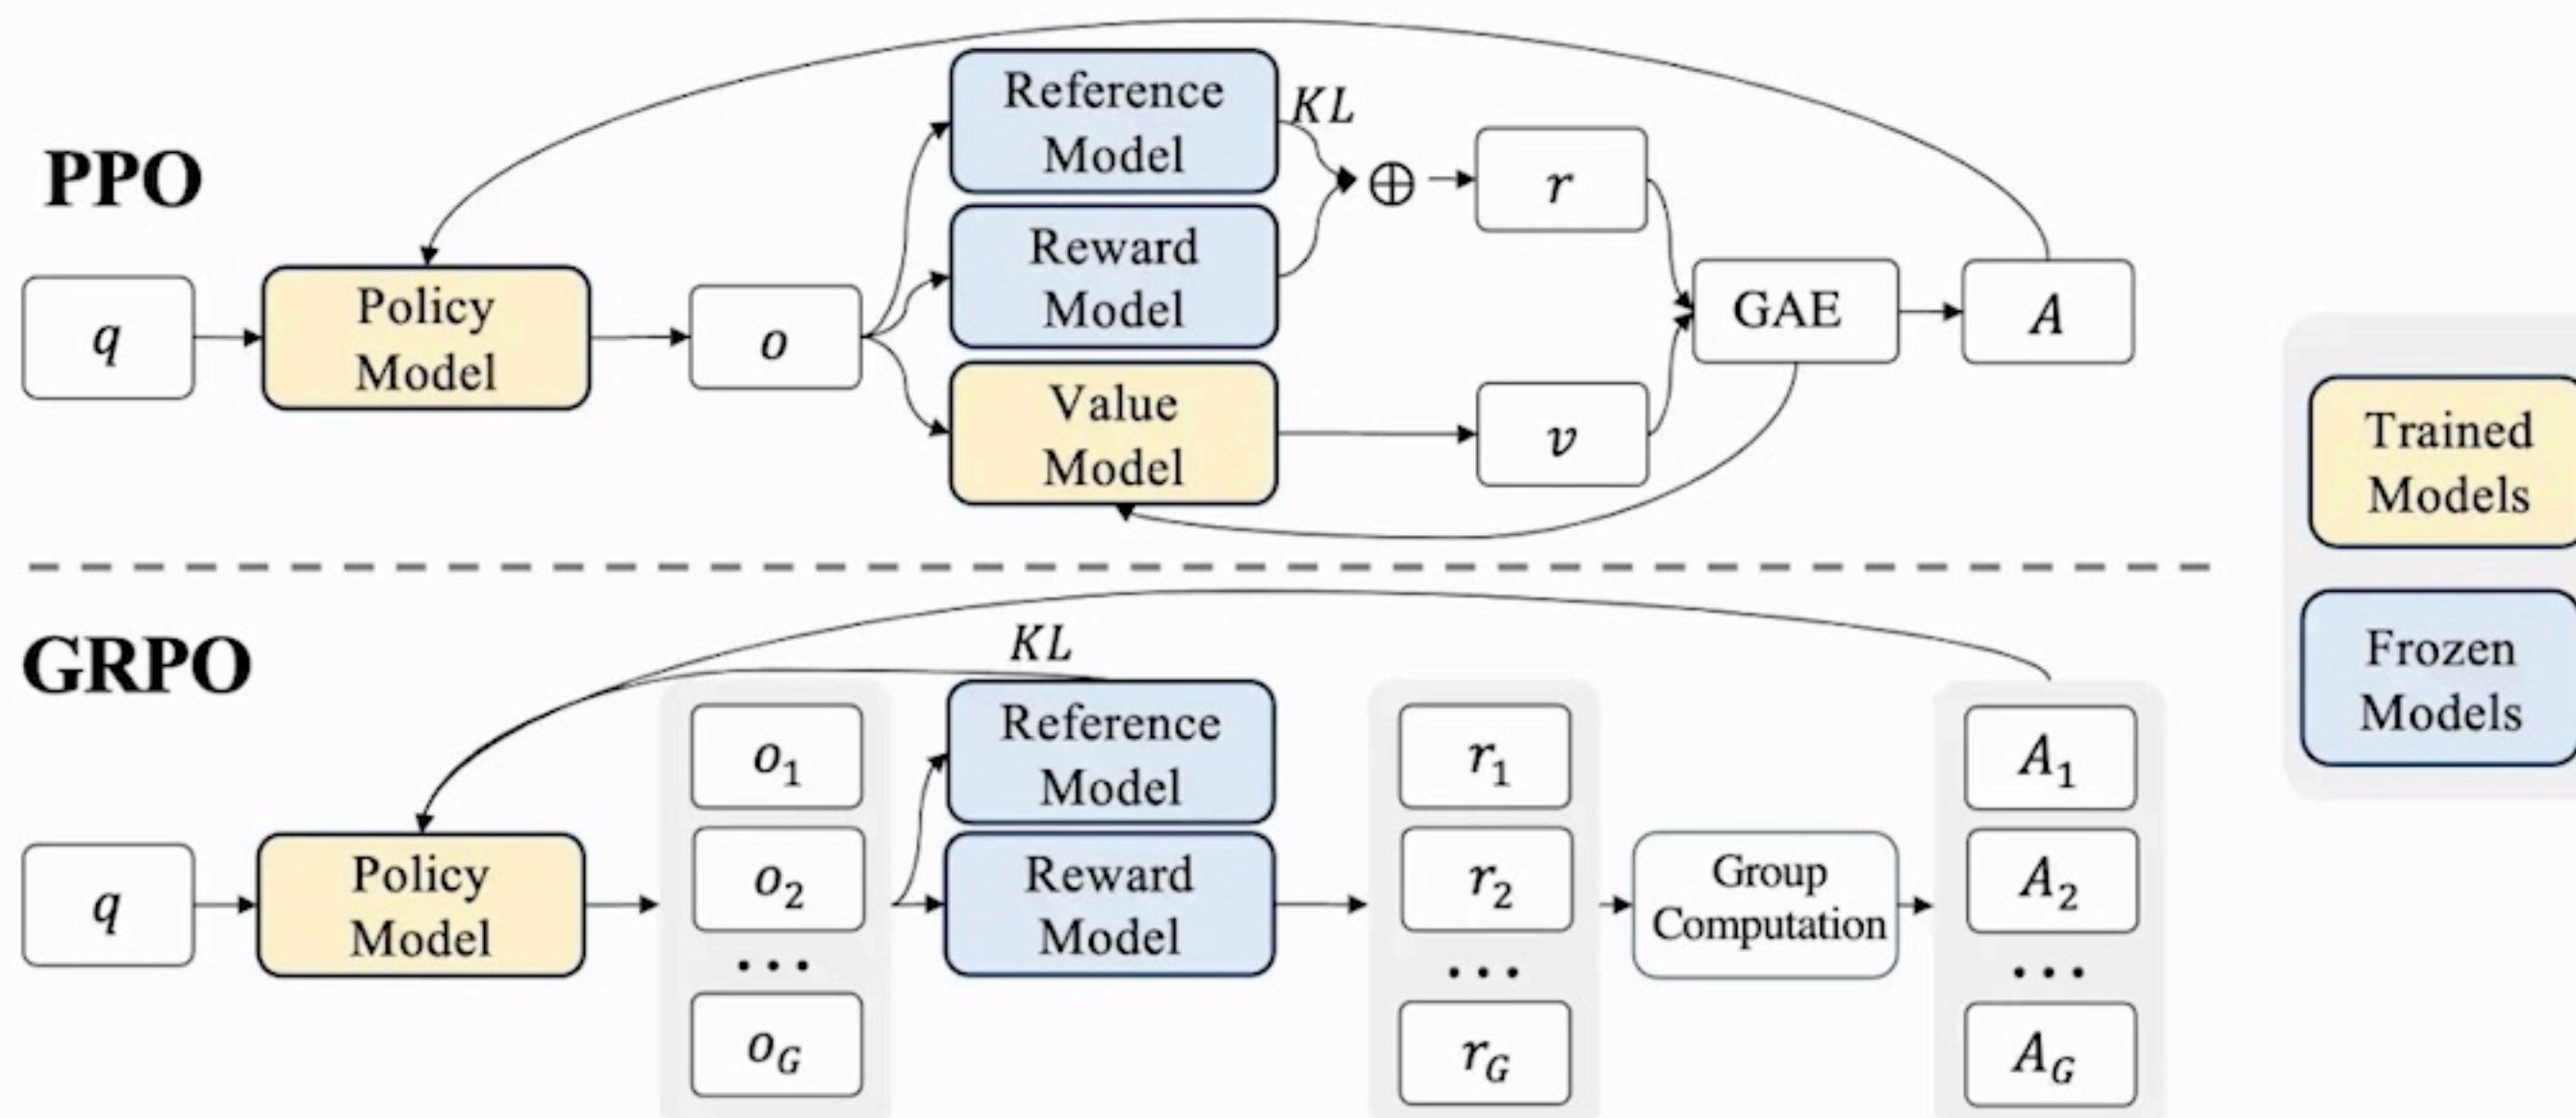
**Response:** import ...

**Test Input 1:** "ABCDABCDBC"

**Test Output 1:** "ABCD"

- Requires preparation of ground truth for math, unit tests for coding, or sandbox execution environment for multi-turn agentic behavior
- More reliable than reward model in those domains
- Used more often for training reasoning models

# Policy Training in Online RL



$$\mathcal{J}_{PPO}(\theta) = \mathbb{E}[q \sim P(Q), o \sim \pi_{\theta_{old}}(O|q)] \frac{1}{|o|} \sum_{t=1}^{|o|} \min \left[ \frac{\pi_{\theta}(o_t|q, o_{<t})}{\pi_{\theta_{old}}(o_t|q, o_{<t})} A_t, \text{clip} \left( \frac{\pi_{\theta}(o_t|q, o_{<t})}{\pi_{\theta_{old}}(o_t|q, o_{<t})}, 1 - \epsilon, 1 + \epsilon \right) A_t \right]$$

Source: Shao, Zhihong, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang et al. "Deepseekmath: Pushing the limits of mathematical reasoning in open language models." arXiv preprint arXiv:2402.03300 (2024).